



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Activation Functions and Convolutional Neural Networks

A. Maier, V. Christlein, K. Breininger, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle,
K. Packhäuser, M. Zinnen

Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 24, 2023



Outline

Activation Functions

Convolutional Neural Networks

Convolutional Layers

Pooling Layers



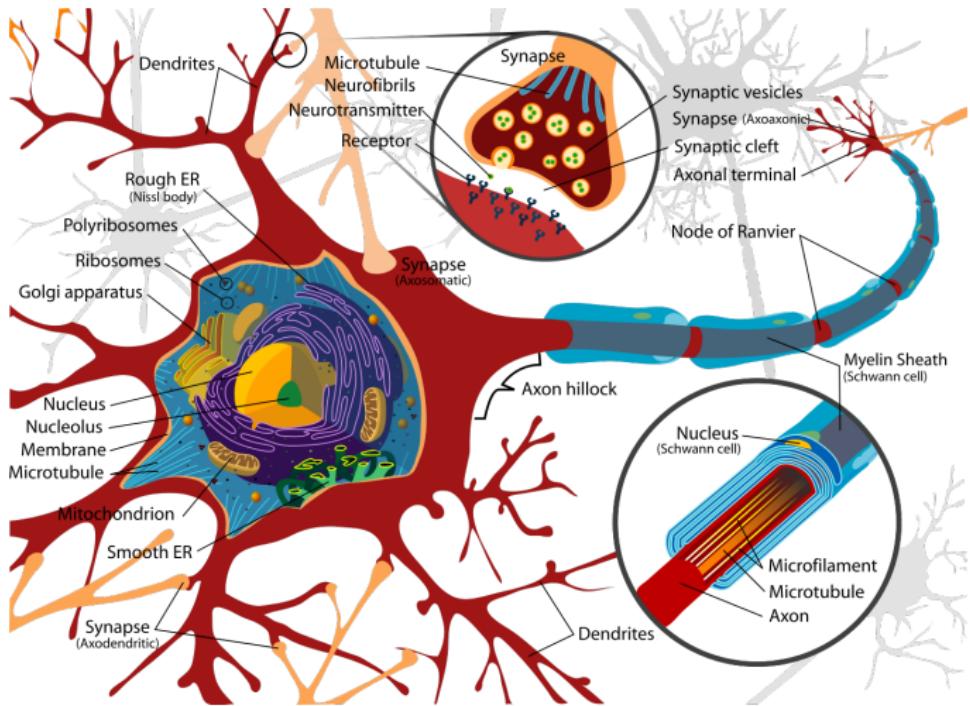
FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

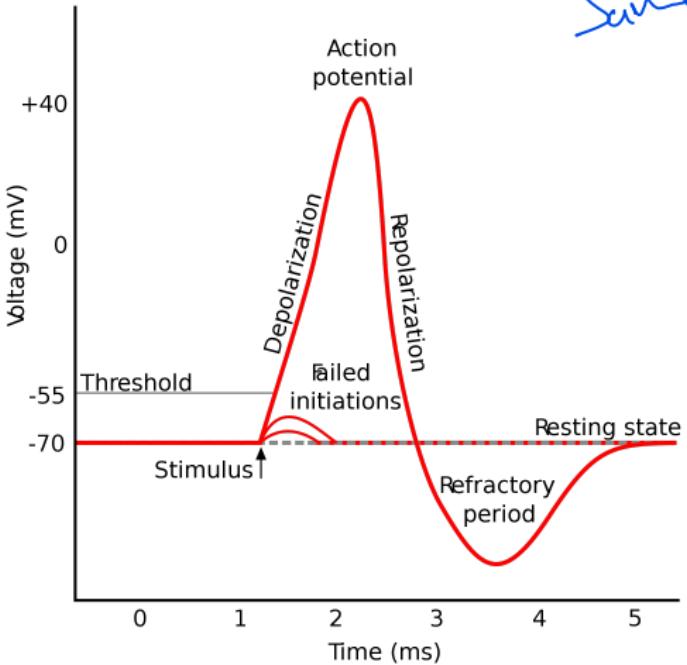
Activation Functions



Biological Activation

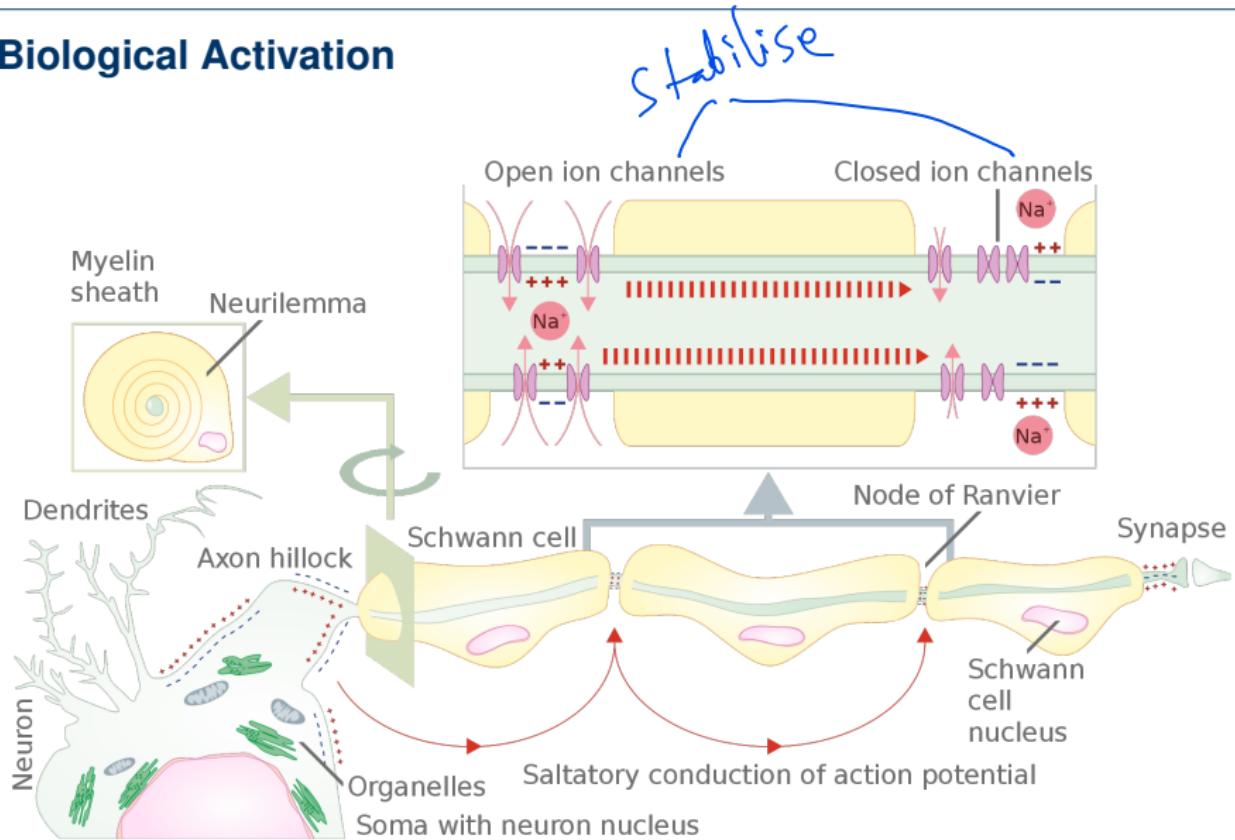


Biological Activation



Always
Same Action
potential

Biological Activation



Summary - Activations in Biological Neurons

- The knowledge lies in the connections
- Both inhibitory and excitatory connections exist
- The synapses anatomically enforce feed forward processing
- However, the connections can be in any direction
- The sum of activations is crucial $\rightarrow \alpha \Sigma > \Theta_{\text{threshold}}$
- Activations are electric spikes
 - With specified intensity
 - Which also encode information over time

whole process
Run at a
frequency.

Activations in Artificial Neural Networks so far

- Non-linear activation functions enable **universal function approximation**:
→ Highly important for a powerful network

Activations in Artificial Neural Networks so far

without non-linearity \rightarrow just matrix multiply

- Non-linear activation functions enable **universal function approximation**:
→ Highly important for a powerful network
- Compared to biology:
 - ✓ Heaviside step (sign) function can model **all or nothing response**
 - ✗ Artificial activations have **no time component** (model by activation strength?)

Activations in Artificial Neural Networks so far

- Non-linear activation functions enable **universal function approximation**:
→ Highly important for a powerful network
- Compared to biology:
 - ✓ Heaviside step (sign) function can model **all or nothing response**
 - ✗ Artificial activations have **no time component** (model by activation strength?)
- Sign function is **mathematically undesirable**:

$$f'(x) = \begin{cases} \infty & \text{for } x = 0 \\ 0 & \text{else} \end{cases}$$



↳ back-propagation

Activations in Artificial Neural Networks so far

- Non-linear activation functions enable **universal function approximation**:
 → Highly important for a powerful network
- Compared to biology:
 - ✓ Heaviside step (sign) function can model **all or nothing response**
 - ✗ Artificial activations have **no time component** (model by activation strength?)
- Sign function is mathematically undesirable:

$$f'(x) = \begin{cases} \infty & \text{for } x = 0 \\ 0 & \text{else} \end{cases}$$

↳ back-propagation

- So far: Sigmoid-function $f(x) = \frac{1}{1+\exp(-x)}$

Activations in Artificial Neural Networks so far

- Non-linear activation functions enable **universal function approximation**.
→ Highly important for a powerful network
- Compared to biology:
 - ✓ Heaviside step (sign) function can model **all or nothing response**
 - ✗ Artificial activations have **no time component** (model by activation strength?)
- Sign function is mathematically undesirable:

$$f'(x) = \begin{cases} \infty & \text{for } x = 0 \\ 0 & \text{else} \end{cases}$$

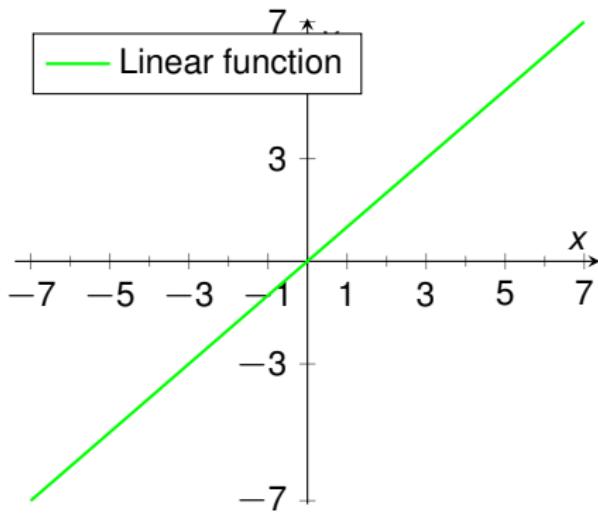
↳ back-propagation

- So far: Sigmoid-function $f(x) = \frac{1}{1+\exp(-x)}$
- Can we do better?

(Analytic derivative is easy)

$$f(u)(1-f(u))$$

Linear Activation Function



Linear function with parameter α :

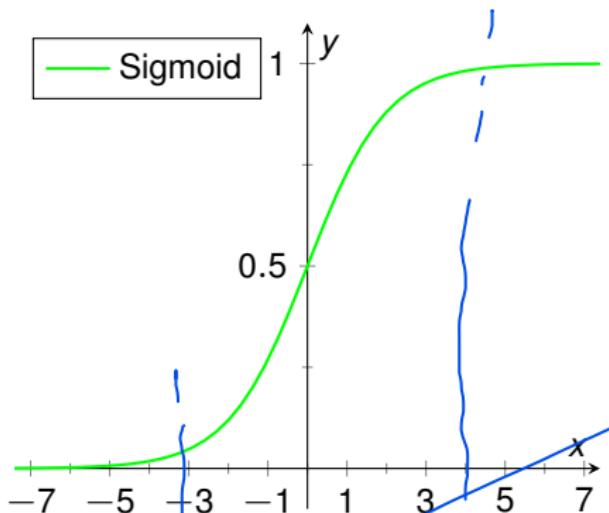
$$f(x) = \alpha x$$

$$f'(x) = \alpha$$

Doesn't allow
for DNN

- + Simple
- Does not introduce non-linearity
- + Therefore, it renders the optimization problem convex
- Listed here mainly for completeness

Sigmoid Activation Function



Sigmoid (logistic function)

$$f(x) = \frac{1}{1 + \exp(-x)}$$

$$f'(x) = f(x)(1 - f(x))$$

f'

- Close to biological model, but differentiable
- + Probabilistic output
- Saturates for $x \ll 0$ and $x \gg 0$
- Not zero-centered

(farther from $(-3, 3)$)
slow (\Rightarrow)
saturation

Zero-Centering

- Sigmoid: $f : \mathbb{R} \mapsto]0, 1[$

Zero-Centering

- Sigmoid: $f : \mathbb{R} \mapsto]0, 1[$
- Output of activation always +
→ ∇_w will either be all + or all -

(Pushing to zero.)

Zero-Centering

Problem

always

values
 $\in [0, 1]$
 positive

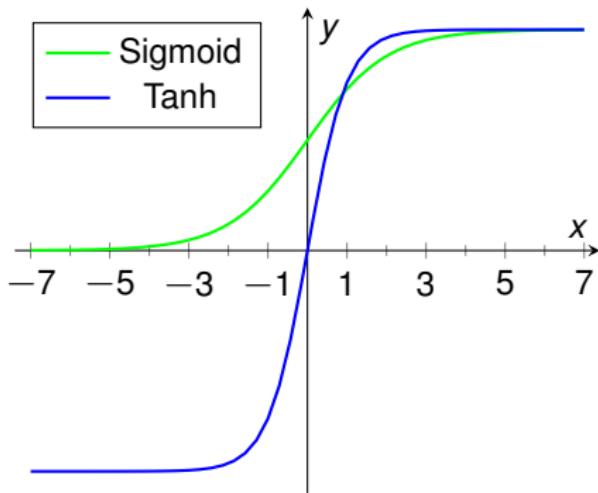
- Sigmoid: $f : \mathbb{R} \mapsto]0, 1[$
- Output of activation always +
 $\rightarrow \nabla_w$ will either be all + or all -
- A mean $\mu = 0$ of the input distribution will always be shifted to $\mu > 0$
 - \rightarrow co-variate shift of successive layers
 - \rightarrow layers constantly have to adapt to the shifting distribution
- Batch learning reduces the variance σ of the updates

Handled.

Cat light
 classify Example

centered around near

Tanh Activation Function



Tanh

$$f(x) = \tanh(x)$$

$$f'(x) = 1 - f(x)^2$$

+ Zero-centered (LeCun '91)

- Shifted version of sigmoid σ : $\tanh(x) = 2\sigma(2x) - 1$

→ Still saturates

→ Still causes vanishing gradients

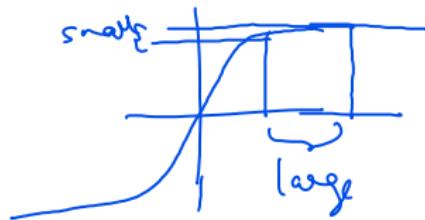
→ Dead ReLU's
 $\{ (0, 0) \}$
 $\{ (2, -2) \}$ - still exists

Vanishing and Exploding Gradients

- Essence of learning: How does x affect y ?
- Sigmoid/tanh map **large regions of X** to a **small range in Y**

Vanishing and Exploding Gradients

- Essence of learning: How does x affect y ?
- Sigmoid/tanh map **large regions** of X to a **small range** in Y
- **Large changes** in $x \mapsto$ minimal changes in y
→ gradient vanishes



Vanishing and Exploding Gradients

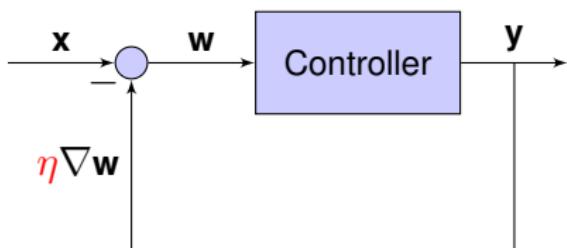
- Essence of learning: How does x affect y ?
- Sigmoid/tanh map **large regions** of X to a **small range** in Y
- **Large changes** in $x \mapsto$ minimal changes in y
 - gradient vanishes
- Problem is amplified by back-propagation: Multiplication of small gradients

Vanishing and Exploding Gradients

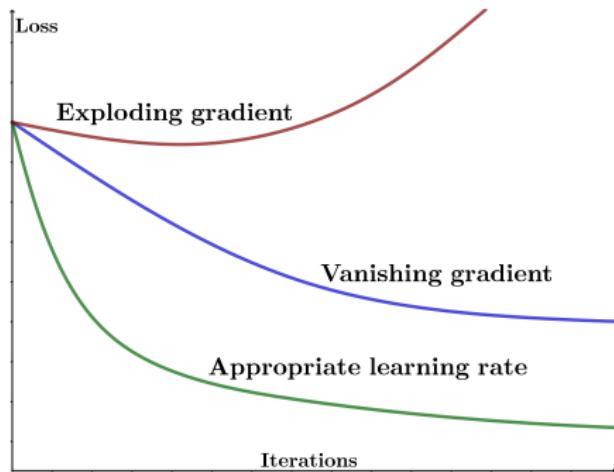
- Essence of learning: How does x affect y ?
 - Sigmoid/tanh map large regions of X to a small range in Y
 - **Large changes** in $x \mapsto$ minimal changes in y
 → gradient vanishes
 - Problem is amplified by back-propagation: Multiplication of small gradients
 - Related problem: Exploding gradients
- Deep layers; faster
 vanishing.

$$[E_{1,1}] \approx y_1.$$

Recap: Feedback Loop - Vanishing and Exploding Gradients



Analogy to control theory



- If η is to high \mapsto **positive feedback** \mapsto loss grows **without bounds**
- If η is to small \mapsto **negative feedback** \mapsto **gradient vanishes**
- Choice of η is **critical** for learning

Can be done
by activation
functions
as well

NEXT TIME
ON DEEP LEARNING



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Activation Functions and Convolutional Neural Networks - Part 2

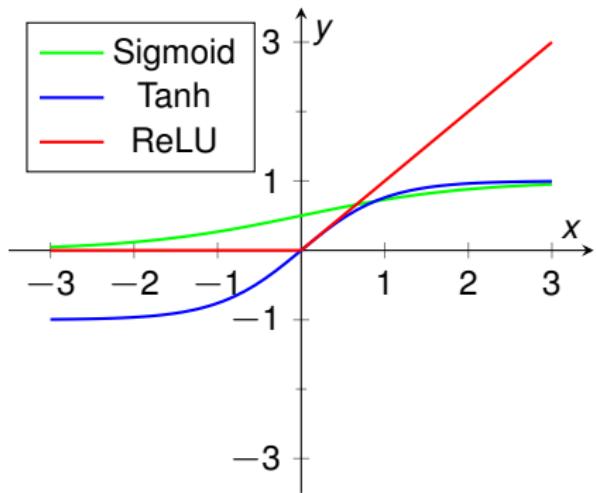
A. Maier, V. Christlein, K. Breininger, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle,
K. Packhäuser, M. Zinnen

Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 24, 2023



Rectified Linear Units (ReLU)



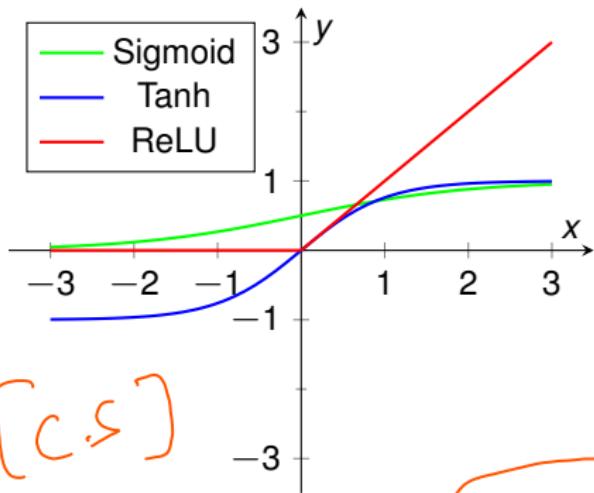
Rectified Linear Unit (ReLU)

$$f(x) = \max(0, x)$$

$$f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{else} \end{cases}$$

- fast
- no vanishing gradient
- linear
- piecewise

Rectified Linear Units (ReLU)



Rectified Linear Unit (ReLU)

$$f(x) = \max(0, x)$$

$$f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{else} \end{cases}$$

- + Good generalization due to piece-wise linearity
- + Speed up during learning ($\approx 6x$ (Krizhevsky '12))
- + No vanishing gradient problem
- Not zero-centered

More on ReLUs

Earlier - more than 3 layers;
Gradient vanishes

- ReLUs were a **big step forward!**
- ReLUs enable training of **deep** supervised neural networks **without unsupervised pre-training**
- First derivative is 1 if the unit is active, second derivative is 0 almost anywhere
 - **No second-order effects**

Dying ReLUs



- Weight/biases trained to yield negative values for **any x**

Dying ReLUs



- Weight/biases trained to yield negative values for **any x**
- ReLU now only performs: $\mathbf{x} \mapsto 0$
- ReLU does not contribute to dividing the feature space

Dying ReLUs



- Weight/biases trained to yield negative values for **any x**
- ReLU now only performs: $\mathbf{x} \mapsto 0$
- ReLU does not contribute to dividing the feature space
- No more updates because $f'(x) \mapsto 0$
→ Our precious ReLU is “dead”

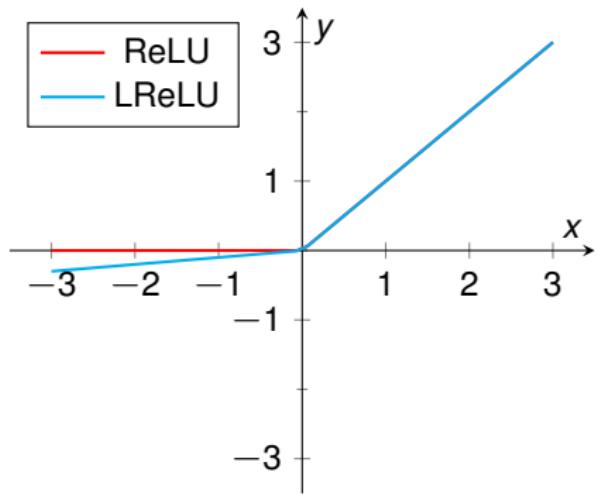
Dying ReLUs



Stops at this point;
b/c no updates

- Weight/biases trained to yield negative values for **any x**
- ReLU now only performs: $x \mapsto 0$
- ReLU does not contribute to dividing the feature space
- No more updates because $f'(x) \mapsto 0$
 - Our precious ReLU is “dead”
- Often related to a (too) high learning rate

Activation Function



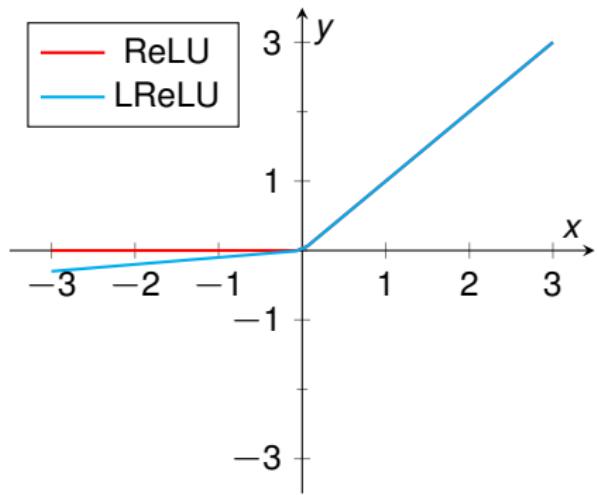
Leaky ReLU / Parametric ReLU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{else} \end{cases}$$

$$f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ \alpha & \text{else} \end{cases}$$

- + Fixes dying ReLU problem
- Leaky ReLU: $\alpha = 0.01$ [5]

Activation Function



Leaky ReLU / Parametric ReLU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{else} \end{cases}$$

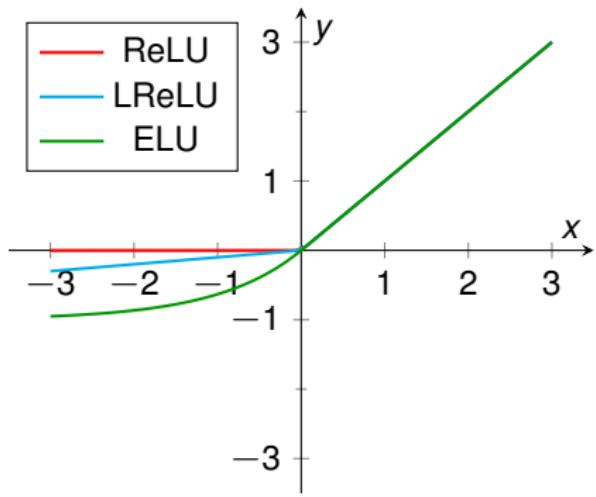
$$f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ \alpha & \text{else} \end{cases}$$

+ Fixes dying ReLU problem

- Leaky ReLU: $\alpha = 0.01$ [5]
- Parametric ReLU (PReLU): learn α [2]

↑ trainable
 ↓ Error → reduces better

Exponential Linear Units (ELU)



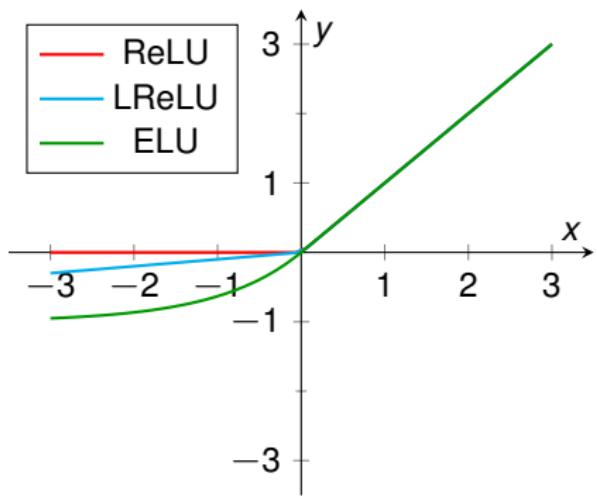
Exponential Linear Unit (ELU)

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(\exp(x) - 1) & \text{else} \end{cases}$$

$$f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ \alpha \exp(x) & \text{else} \end{cases}$$

Smooth function on
negative Halfspace

Exponential Linear Units (ELU)



Exponential Linear Unit (ELU)

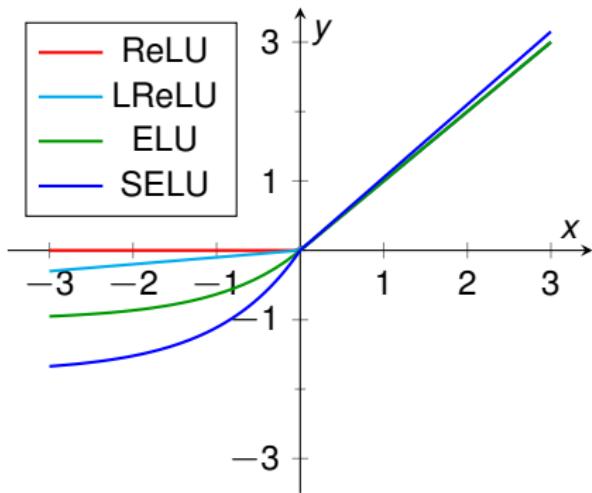
$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(\exp(x) - 1) & \text{else} \end{cases}$$

$$f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ \alpha \exp(x) & \text{else} \end{cases}$$

- + Also no vanishing gradient
- + Reduces shift in activations

also negative α 's can be present

Scaled ELU (SELU)



Scaled Exponential Linear Unit

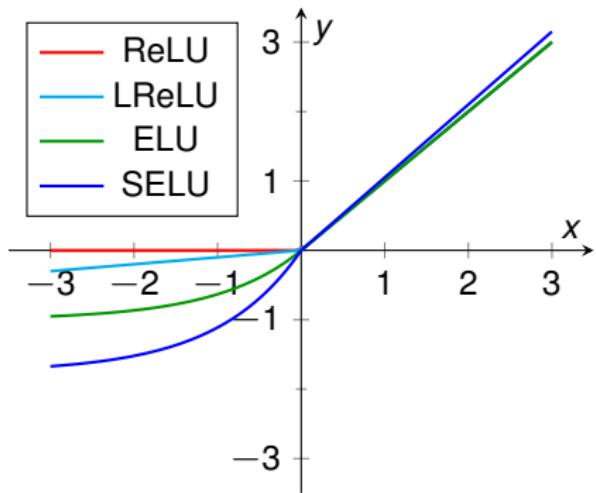
$$f(x) = \lambda \begin{cases} x & \text{if } x > 0 \\ \alpha(\exp(x) - 1) & \text{else} \end{cases}$$

$$f'(x) = \lambda \begin{cases} 1 & \text{if } x > 0 \\ \alpha \exp(x) & \text{else} \end{cases}$$

$$\lambda_{01} = 1.0507$$

$$\alpha_{01} = 1.6733$$

Scaled ELU (SELU)



$$\mu = 0; \sigma = 1$$

Scaled Exponential Linear Unit

$$f(x) = \lambda \begin{cases} x & \text{if } x > 0 \\ \alpha(\exp(x) - 1) & \text{else} \end{cases}$$

$$f'(x) = \lambda \begin{cases} 1 & \text{if } x > 0 \\ \alpha \exp(x) & \text{else} \end{cases}$$

$$\lambda_{01} = 1.0507$$

$$\alpha_{01} = 1.6733$$

use this
if \circ mean
 \approx variance

No problem with
internal Covariate
shift

- Alternative variant of ReLU [3]
- Idea: Self-normalizing - $\mu = 0, \sigma = 1 \mapsto \lambda_{01}, \alpha_{01}$
- Alternative to Batch Normalization? (see next lecture)

Other Activation Functions

Other Activation Functions

- Maxout: Learns the activation function [1]

Other Activation Functions

- Maxout: Learns the activation function [1]
- Radial basis functions

Other Activation Functions

- Maxout: Learns the activation function [1]
- Radial basis functions
- Softplus $f(x) = \ln(1 + e^x)$: less efficient than ReLU

Other Activation Functions

- Maxout: Learns the activation function [1]
- Radial basis functions
- Softplus $f(x) = \ln(1 + e^x)$: less efficient than ReLU

This is getting ridiculous - what should we use?

Finding the Optimal Activation Function

Finding the Optimal Activation Function

- Reinforcement learning/search problem

Finding the Optimal Activation Function

- Reinforcement learning/search problem
 - Unfortunately a single “step” means **training a network** from scratch

Finding the Optimal Activation Function

- Reinforcement learning/search problem
 - Unfortunately a single “step” means **training a network** from scratch
 - If you have a **cloud/supercomputer** you can do something like this
- “Searching for Activation Functions” [6] published by Google on Oct 16, 2017

Finding the Optimal Activation Function

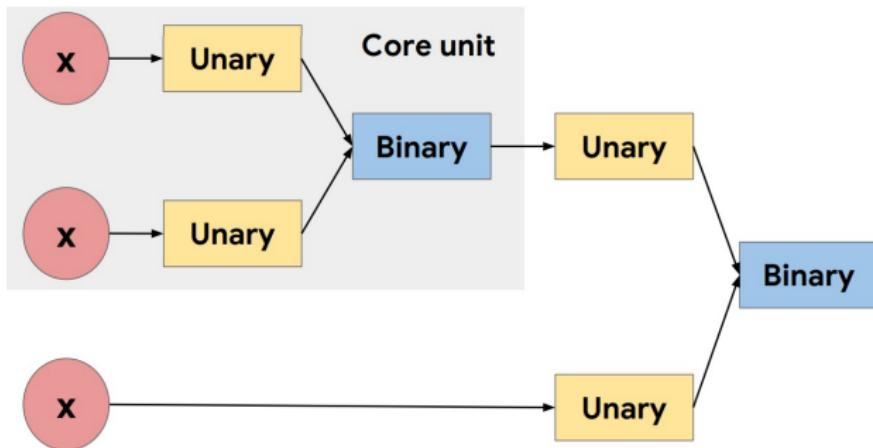
later

- Reinforcement learning/search problem
 - Unfortunately a single “step” means **training a network** from scratch
 - If you have a **cloud/supercomputer** you can do something like this
- “Searching for Activation Functions” [6] published by Google on Oct 16, 2017

Strategy

1. Define a search-space
2. Perform the search using a RNN with reinforcement learning
3. Use the best result

Searching for Activation Functions [6]

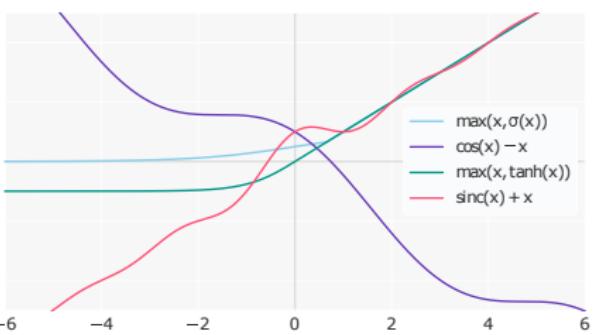
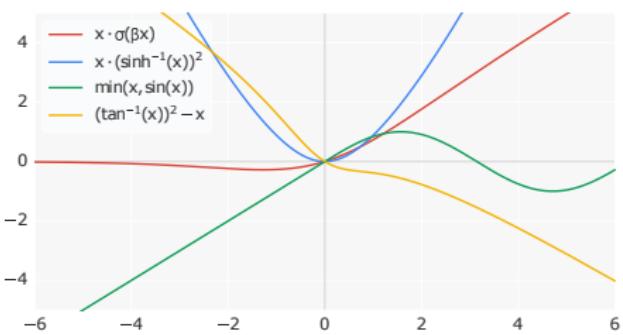


Search space for activation functions

Search Space

Source: <https://arxiv.org/pdf/1710.05941.pdf>

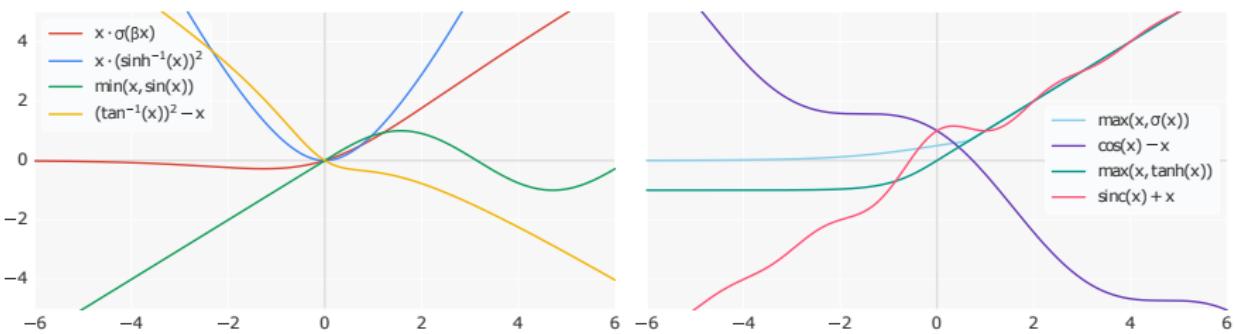
Searching for Activation Functions [6]



- We don't have a **cloud**, but we can use those

Source: <https://arxiv.org/pdf/1710.05941.pdf>

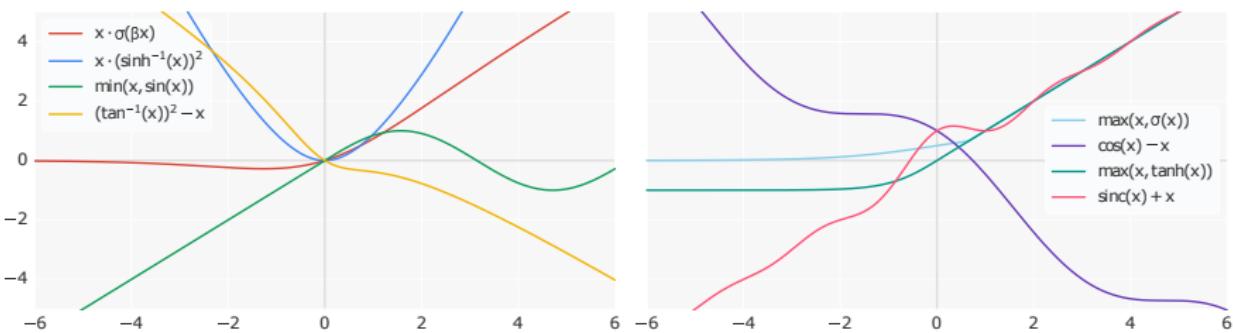
Searching for Activation Functions [6]



- We don't have a **cloud**, but we can use those
- **Complicated** activation functions **didn't perform well**

Source: <https://arxiv.org/pdf/1710.05941.pdf>

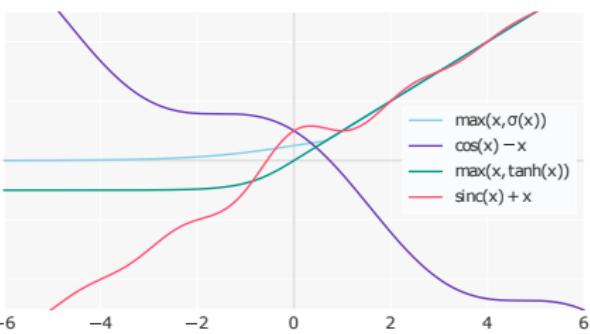
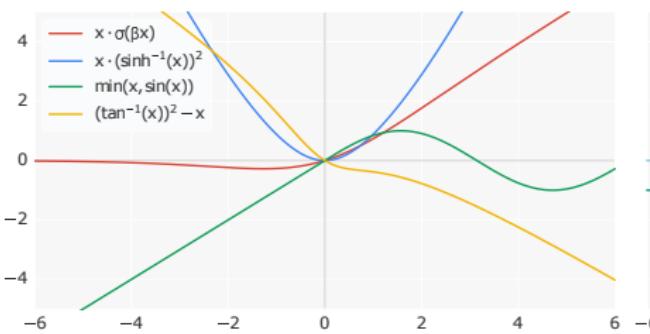
Searching for Activation Functions [6]



- We don't have a **cloud**, but we can use those
- **Complicated** activation functions **didn't perform well**
- They now call $x \cdot \sigma(\beta x)$ the “Swish” function

Source: <https://arxiv.org/pdf/1710.05941.pdf>

Searching for Activation Functions [6]



- We don't have a **cloud**, but we can use those
- **Complicated** activation functions **didn't perform well**
- They now call $x \cdot \sigma(\beta x)$ the "Swish" function
- Has been proposed before as "Sigmoid-weighted Linear Unit" [7]

Source: <https://arxiv.org/pdf/1710.05941.pdf>

Let's look into their results in detail

Disclaimer

Never show tables in your slides. Try hard to find a better representation.

Source: <https://arxiv.org/pdf/1710.05941.pdf>

Let's look into their results in detail

Model	Top-1 Acc. (%)		
LReLU	79.5	79.5	79.6
PReLU	79.7	79.8	80.1
Softplus	80.1	80.2	80.4
ELU	75.8	79.9	80.0
SELU	79.0	79.2	79.2
GELU	79.6	79.6	79.9
ReLU	79.5	79.6	79.8
Swish-1	80.2	80.3	80.4
Swish	80.2	80.2	80.3

Inception-Resnet-V2 architecture trained on ImageNet.

Source: <https://arxiv.org/pdf/1710.05941.pdf>

Let's look into their results in detail

Model	Top-1 Acc. (%)		
LReLU	79.5	79.5	79.6
PReLU	79.7	79.8	80.1
Softplus	80.1	80.2	80.4
ELU	75.8	79.9	80.0
SELU	79.0	79.2	79.2
GELU	79.6	79.6	79.9
ReLU	79.5	79.6	79.8
Swish-1	80.2	80.3	80.4
Swish	80.2	80.2	80.3

- Are any of these differences significant?

Not much

is observation
a result of
randomness

Inception-Resnet-V2 architecture trained on ImageNet.

Source: <https://arxiv.org/pdf/1710.05941.pdf>

Summary Activation Functions

- Go-to solution: **ReLU**

Summary Activation Functions

- Go-to solution: **ReLU**
- SELUs have **attractive** theoretical properties.

Summary Activation Functions

- Go-to solution: **ReLU**
- SELUs have **attractive** theoretical properties.
- ... but try ReLU first.

Summary Activation Functions

- Go-to solution: **ReLU**
- SELUs have **attractive** theoretical properties.
- ... but try ReLU first.
- Best activation function is a **difficult, expensive optimization problem**.

Summary Activation Functions

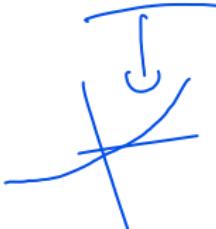
- Go-to solution: **ReLU**
- SELUs have **attractive** theoretical properties.
- ... but try ReLU first.
- Best activation function is a **difficult, expensive optimization problem**.

Summary Activation Functions

- Go-to solution: **ReLU**
- SELUs have **attractive** theoretical properties.
- ... but try ReLU first.
- Best activation function is a **difficult, expensive optimization problem**.

What we know about good activation functions:

- They have almost linear areas to prevent **vanishing gradients**.
- They have **saturating areas** to provide **non-linearity**.
- They should be **monotonic**.



where gradient
 $= 0$

linearity
 \Downarrow
 no vanishing
 gradient

NEXT TIME
ON DEEP LEARNING

Activation Functions and Convolutional Neural Networks - Part 3

A. Maier, V. Christlein, K. Breininger, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle,
K. Packhäuser, M. Zinnen

Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 24, 2023





FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

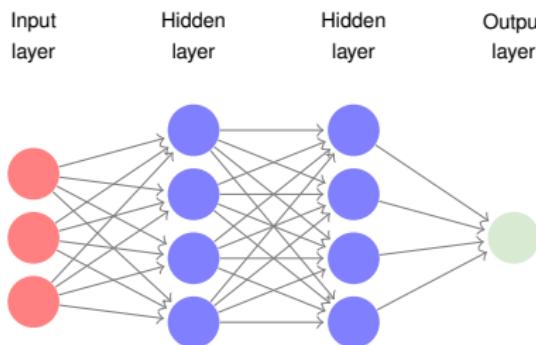
Convolutional Neural Networks

↳ Powerful building blocks



Motivation

- So far: Fully connected layers - each input is connected to each node
- Very powerful: Can represent any kind of (linear) relationship between inputs



Matrix multiplications

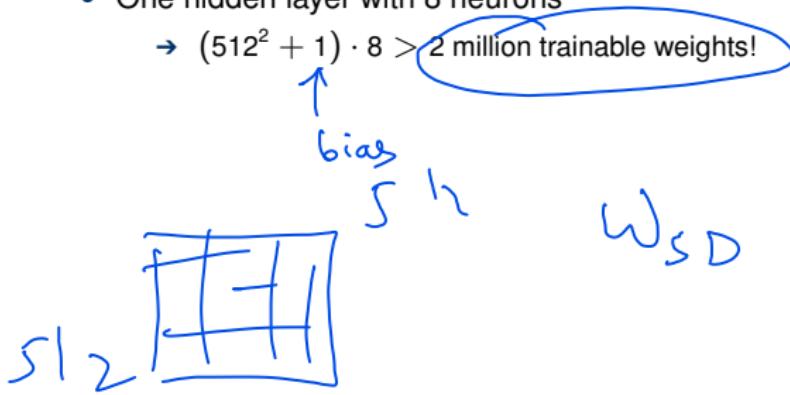
Motivation

- So far: Fully connected layers - each input is connected to each node
- Very powerful: Can represent any kind of (linear) relationship between inputs
- Large part of machine learning: images/videos/sounds

Motivation

- So far: Fully connected layers - each input is connected to each node
- Very powerful: Can represent any kind of (linear) relationship between inputs
- Large part of machine learning: images/videos/sounds
- Assume we have:
 - An image with size 512×512 pixels
 - One hidden layer with 8 neurons

$$\rightarrow (512^2 + 1) \cdot 8 > 2 \text{ million trainable weights!}$$



Motivation (cont.)

So the size is a problem. Is there something else?

Source: <https://news.nationalgeographic.com>

Motivation (cont.)

So the size is a problem. Is there something else?

- Example: Classify between cat and dog

→ not the only problem



Source: <https://news.nationalgeographic.com>

Motivation (cont.)

So the size is a problem. Is there something else?

- Example: Classify between cat and dog



Source: <https://news.nationalgeographic.com>

Motivation (cont.)

So the size is a problem. Is there something else?

- Example: Classify between cat and dog
- Pixels are **bad** features!
 - Highly correlated
 - Scale dependent
 - Intensity variations
 - ...



Source: <https://news.nationalgeographic.com>

Motivation (cont.)

So the size is a problem. Is there something else?

- Example: Classify between cat and dog
- Pixels are **bad** features!
 - Highly correlated
 - Scale dependent
 - Intensity variations
 - ...
- Pixels are a bad representation from a machine learning point of view

light,
camera

above
reasons

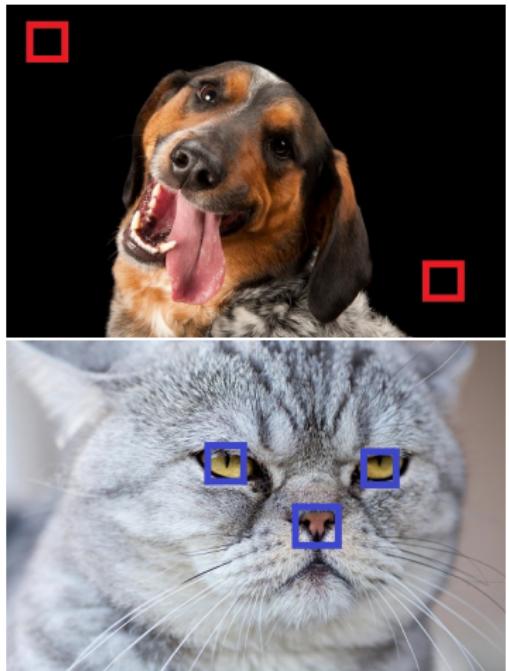
shadow, reflections
noise vs feature



Source: <https://news.nationalgeographic.com>

Motivation (cont.)

Can we find a better representation?



Source: <https://news.nationalgeographic.com>

Motivation (cont.)

Can we find a better representation?

- We have a certain degree of the locality in an image



Source: <https://news.nationalgeographic.com>

Motivation (cont.)

Can we find a better representation?

- We have a certain degree of the locality in an image
- We can find the same “macro features” at different locations

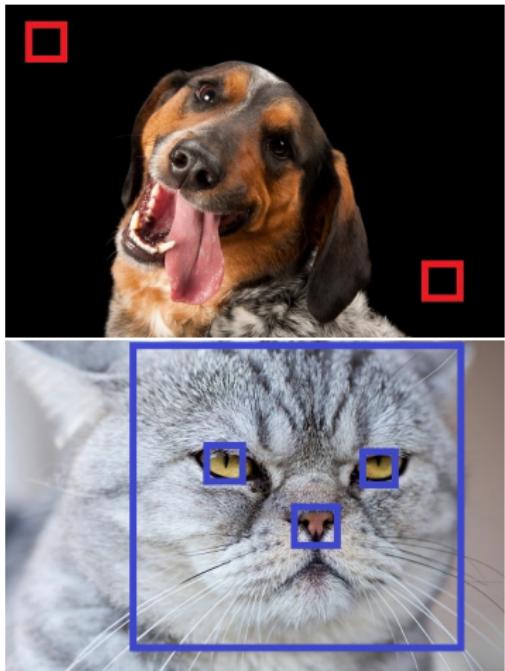


Source: <https://news.nationalgeographic.com>

Motivation (cont.)

Can we find a better representation?

- We have a certain degree of the locality in an image
- We can find the same “macro features” at different locations
- Hierarchy of features:
 - edges + corners \mapsto eyes
 - eyes + nose + ears \mapsto face
 - face + body + legs \mapsto animal



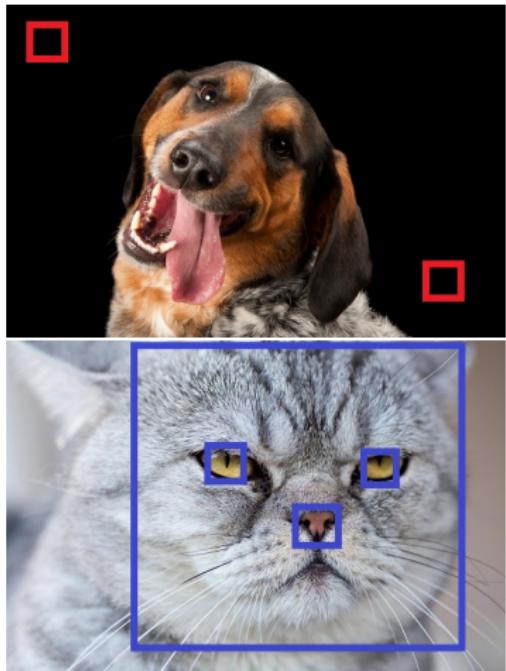
Source: <https://news.nationalgeographic.com>

Motivation (cont.)

[C-S]

Can we find a better representation?

- We have a certain degree of the locality in an image
 - We can find the same “macro features” at different locations
 - Hierarchy of features:
 - edges + corners \mapsto eyes
 - eyes + nose + ears \mapsto face
 - face + body + legs \mapsto animal
 - Composition matters!
 - Learn better representation, then classify!
- low to High*



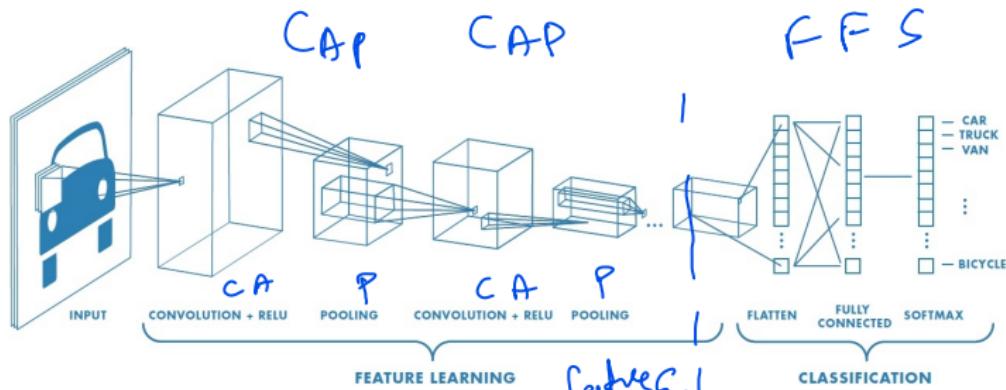
Source: <https://news.nationalgeographic.com>

Convolutional Neural Networks



- Local connectivity → filters
- Use same filters over the whole image → translational equivariance
- Hierarchy of filters working on different scales
- + learning = Convolutional Neural Networks

Convolutional Neural Networks - Architecture



Four essential building blocks:

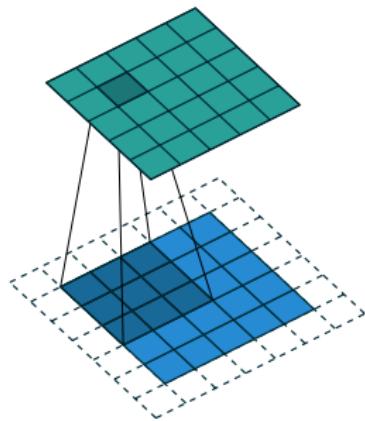
- Convolutional layer: Feature extraction
- Activation function: Nonlinearity
- Pooling layer: Compress and aggregate information, save parameters
- Last layer: Fully-connected for classification

Source: <https://de.mathworks.com/discovery/convolutional-neural-network.html>

Convolutional Layers

Convolutional Layer - Local Connectivity

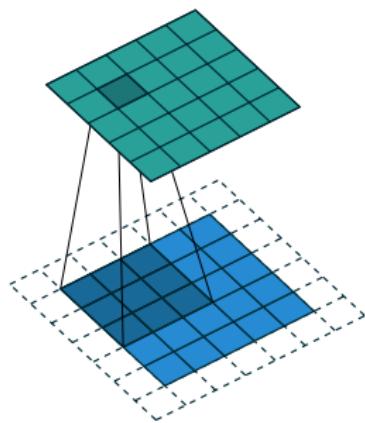
- Exploit spacial structure by only connecting pixels in a neighborhood
- Can be expressed as fully connected layer:
Except for local connections, each entry in W is 0



Source: https://github.com/vdumoulin/conv_arithmetic

Convolutional Layer - Local Connectivity

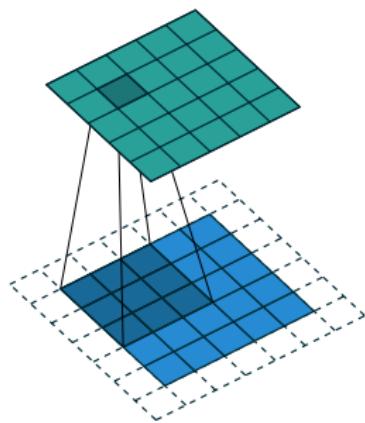
- Exploit spacial structure by only connecting pixels in a neighborhood
- Can be expressed as fully connected layer:
Except for local connections, each entry in \mathbf{W} is 0
- Effective weights: Filter of size 3×3 , 5×5 , 7×7 , ...



Source: https://github.com/vdumoulin/conv_arithmetic

Convolutional Layer - Local Connectivity

- Exploit spatial structure by only connecting pixels in a neighborhood
- Can be expressed as fully connected layer:
Except for local connections, each entry in \mathbf{W} is 0
- Effective weights: Filter of size 3×3 , 5×5 , 7×7 , ...
- Features that are important at one location are likely important anywhere in the image
 - Use the same weights all over: **tied weights** (also **shared weights**).

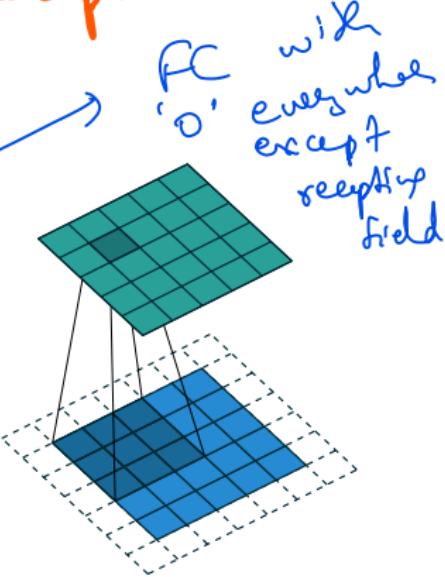


Source: https://github.com/vdumoulin/conv_arithmetic

$FC = ConvL + 'O' \text{ outside receptive field}$

Convolutional Layer - Local Connectivity

- Exploit spacial structure by only connecting pixels in a neighborhood
- Can be expressed as fully connected layer:
Except for local connections, each entry in W is 0
- Effective weights: Filter of size 3×3 , 5×5 , 7×7 , ...
- Features that are important at one location are likely important anywhere in the image
 - Use the same weights all over: tied weights (also shared weights)
 - Convolution with trainable filters



(Modelling
Convolution)

finding best filters

Source: https://github.com/vdumoulin/conv_arithmetic

Recap: Convolution

Convolution

Source: https://github.com/vdumoulin/conv_arithmetic

Recap: Convolution

- Convolution:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(\tau)g(x - \tau)d\tau$$

Recap: Convolution

- Convolution:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(\tau)g(x - \tau)d\tau$$

- Cross-correlation:

$$(f \star g)(x) = \int_{-\infty}^{\infty} f(\tau)g(x + \tau)d\tau$$

Recap: Convolution

- Convolution:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(\tau)g(x - \tau)d\tau$$

- Cross-correlation:

$$(f \star g)(x) = \int_{-\infty}^{\infty} f(\tau)g(x + \tau)d\tau$$

- Cross-correlation is convolution with a flipped kernel g – and vice versa!

$$C = \begin{array}{|c|c|c|} \hline 1 & 6 & 1 \\ \hline 5 & 2 & 1 \\ \hline 5 & 1 & 6 \\ \hline \end{array} \quad | \quad \begin{array}{|c|c|c|} \hline 3 & 2 & 1 \\ \hline 5 & 1 & 6 \\ \hline 1 & 5 & 6 \\ \hline \end{array}$$

Recap: Convolution

- Convolution:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(\tau)g(x - \tau)d\tau$$

different domain
(move negative)



- Cross-correlation:

$$(f \star g)(x) = \int_{-\infty}^{\infty} f(\tau)g(x + \tau)d\tau$$

(move positive)



- Cross-correlation is convolution with a flipped kernel g – and vice versa!
- Implementation: Cross-correlation is frequently used in the forward pass - the weights are initialized randomly anyway

Doesn't matter in training:
it's just a sign

Padding

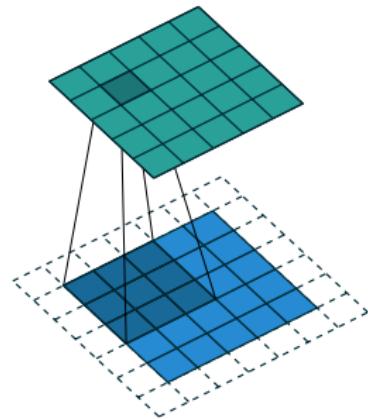
$$[(w - k + 2p) + 1]$$

w

k

p

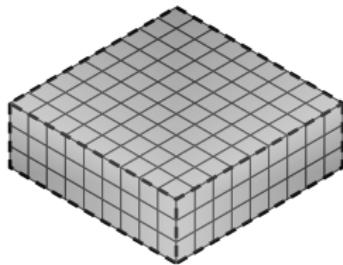
- Convolution reduces image size by $2 \cdot \lfloor \frac{n}{2} \rfloor$ pixels (n: kernel size).
- Necessary to pay attention to the borders:
- 'Same' padding (usually zero padding):
 - Input and output have the same size
- 'valid'/no padding:
 - The output is smaller than the input



Source: https://github.com/vdumoulin/conv_arithmetic

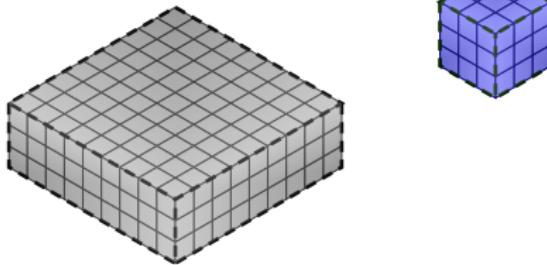
Forward Pass: Multi-channel convolution

- Input of size $X \times Y \times S$, where S is the number of input channels



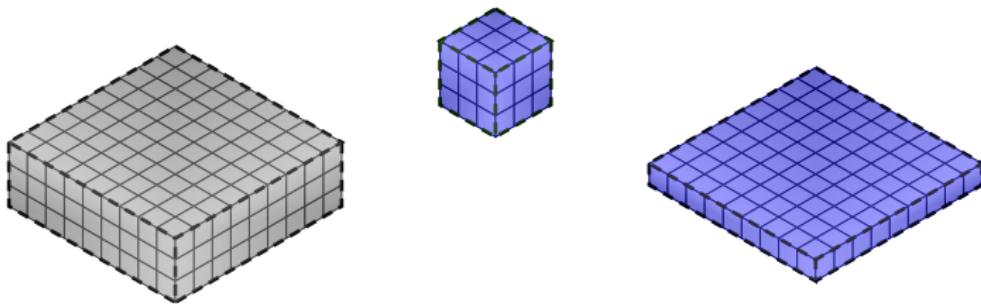
Forward Pass: Multi-channel convolution

- Input of size $X \times Y \times S$, where S is the number of input channels
- H Filters with size $M \times N \times S \mapsto$ **fully connected** across channels



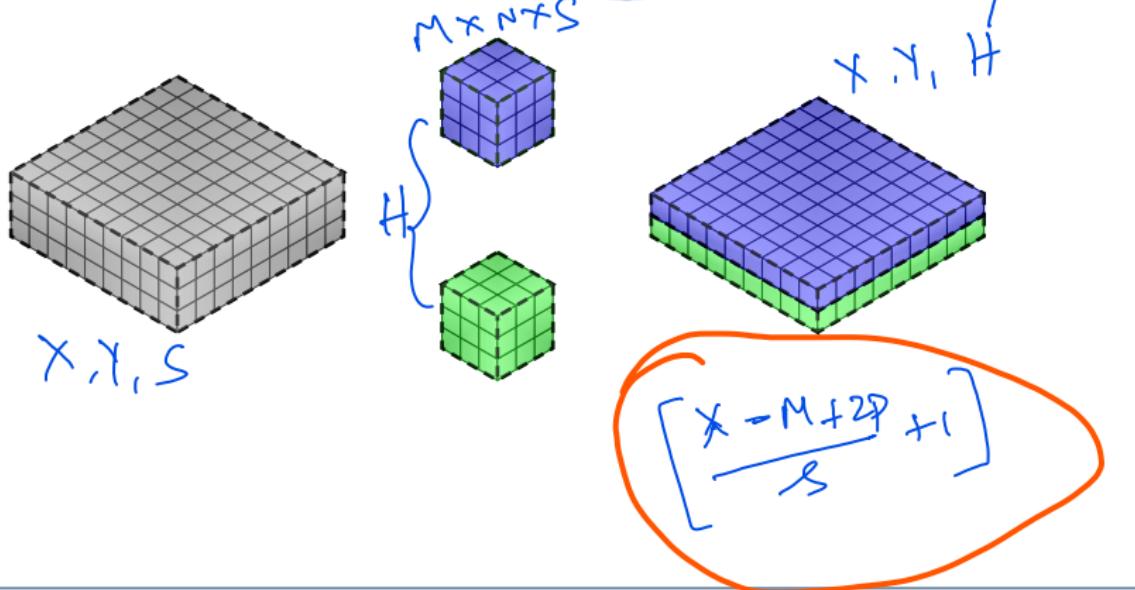
Forward Pass: Multi-channel convolution

- Input of size $X \times Y \times S$, where S is the number of input channels
- H Filters with size $M \times N \times S \mapsto$ **fully connected** across channels
- Output dimensions: $X \times Y \times H$ (with ‘same’ padding)



Forward Pass: Multi-channel convolution

- Input of size $X \times Y \times S$, where S is the number of input channels
- H Filters with size $M \times N \times S$ \mapsto **fully connected** across channels
- Output dimensions: $X \times Y \times H$ (with 'same' padding)



Backward pass: Multi-channel convolution

- Convolution can be expressed as matrix multiplication with matrix W^T using a **Toeplitz matrix** \rightarrow **Circular Matrix (weight sharing)**
- We can use the **same formulas** as for the fully connected layer!

$$\begin{aligned} y &= W^T X \\ \nabla W &= E_{l-1} X^T \\ E_{l-1} &= W \cdot E_l \end{aligned}$$

$$\begin{aligned} E_{l-1} &= W^T E_l \\ \nabla W &= E_l X^T \end{aligned} \quad \left. \begin{array}{l} \text{Same Update} \\ \text{formulas} \end{array} \right\}$$

where

- X is the input and
- E_l/E_{l-1} is the error in layer $l/l - 1$

(ES)

$$E_l = \frac{\partial L}{\partial y}$$

Backward pass: Multi-channel convolution

- Convolution can be expressed as matrix multiplication with matrix \mathbf{W} : using a **Toeplitz matrix**
- We can use the **same formulas** as for the fully connected layer!

$$\mathbf{E}_{l-1} = \mathbf{W}^T \mathbf{E}_l$$

$$\nabla \mathbf{W} = \mathbf{E}_l \mathbf{X}^T$$

where

- \mathbf{X} is the input and
- $\mathbf{E}_l / \mathbf{E}_{l-1}$ is the error in layer $l // l - 1$
- Backward pass can also be expressed as convolutions ↪ exercise

Convolutional Layer - What have we gained?

- Stack multiple filters to get a trainable filter bank.
- Layer with 8 filters (nodes) with 5×5 neighborhood
 - $5^2 \cdot 8 = \underline{200}$ weights

Convolutional Layer - What have we gained?

- Stack multiple filters to get a trainable filter bank.
- Layer with 8 filters (nodes) with 5×5 neighborhood
 $\rightarrow 5^2 \cdot 8 = \underline{200}$ weights
- Convolution **Independent** of image size!
- Much more training data for one weight!

$m \times N \times H$
 $S + S + 8$
 diff from 2^M
 earlier

Strided Convolutions

- Instead of multiplying the filter at each pixel position, we can **skip** some **positions**
- **Stride** s describes the offset



Source: https://en.wikipedia.org/wiki/Monty_Python%27s_Flying_Circus, Dinsdale

Strided Convolutions

- Instead of multiplying the filter at each pixel position, we can **skip** some **positions**
- **Stride** s describes the offset
- **Reduces** the **size** of the output by a factor of s



Source: https://en.wikipedia.org/wiki/Monty_Python%27s_Flying_Circus_episode,_Dinsdale

Strided Convolutions

- Instead of multiplying the filter at each pixel position, we can **skip some positions**
- Stride s** describes the offset
- Reduces the size of the output by a factor of s**
- Mathematically: Convolution + subsampling**

$\text{stride} = \text{Conv} + \text{Subsampling}$

$(C \cdot S)$



Source: https://en.wikipedia.org/wiki/Monty_Python%27s_Flying_Circus_episode,_Dinsdale

Strided Convolutions

Strided Convolution with stride $s = 2$

Source: https://github.com/vdumoulin/conv_arithmetic

Dilated/Atrous Convolutions

- Additional variant of convolution in neural networks
- Dilate convolution kernel: Skip certain pixels → i/p
- Goal: Wider receptive field with less parameters/weights



Receptive field
is
not
Connected

Dilated Convolution

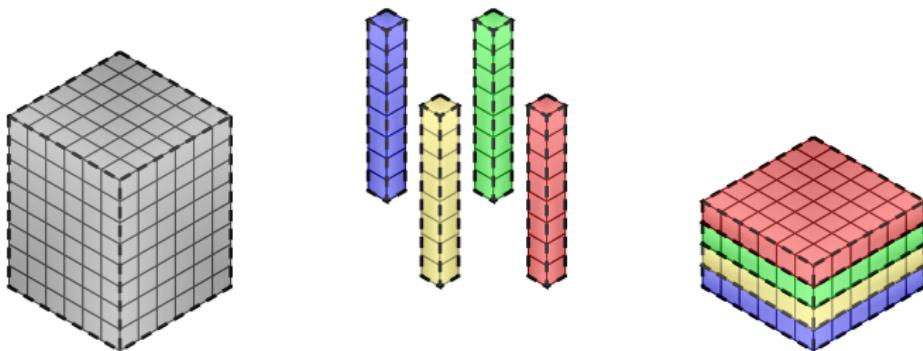
Source: https://github.com/vdumoulin/conv_arithmetic

1 × 1 Convolution Concept

- So far: H filters with neighborhood $3 \times 3, 5 \times 5, \dots$ and ‘depth’ S
- Filters are fully connected in ‘depth’ direction

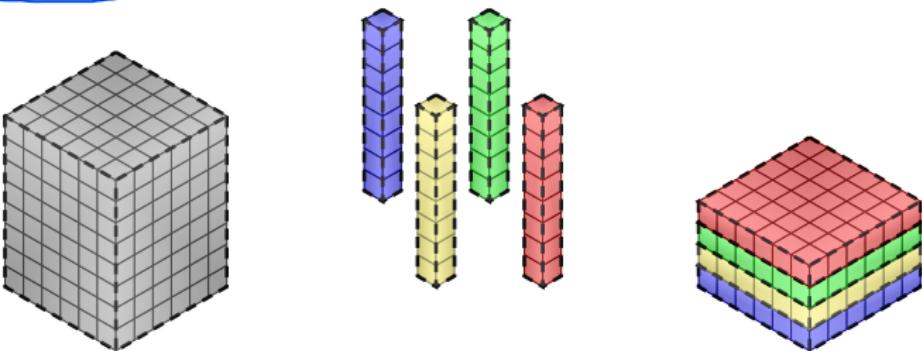
1×1 Convolution Concept

- So far: H filters with neighborhood $3 \times 3, 5 \times 5, \dots$ and ‘depth’ S
- Filters are fully connected in ‘depth’ direction
- We can decrease the neighborhood to 1×1
- And **just** use the fully connected property in the depth dimension



1 × 1 Convolution Concept

- So far: H filters with neighborhood $3 \times 3, 5 \times 5, \dots$ and 'depth' S
- Filters are fully connected in 'depth' direction
- We can decrease the neighborhood to 1×1
- And just use the fully connected property in the depth dimension



- Dimensionality reduction/expansion from S channels to H channels
- If we flatten the input, 1 × 1 convolutions are fully connected layer!

1×1 Convolution Concept (cont.)

- First described in “Network in Network” by Lin et al. [4]
- 1×1 convolutions simply calculate **inner products** at each position
- Simple and efficient method to **decrease** the **size** of a network
- **Learns** dimensionality reduction, e.g., can reduce redundancy in your feature maps

1×1 Convolution Concept (cont.)

- First described in “Network in Network” by Lin et al. [4]
 - 1×1 convolutions simply calculate **inner products** at each position
 - Simple and efficient method to **decrease the size** of a network
 - **Learns** dimensionality reduction, e.g., can reduce redundancy in your feature maps
-
- Equivalent but more flexible: $N \times N$ convolution

**NEXT TIME
ON DEEP LEARNING**



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Activation Functions and Convolutional Neural Networks - Part 4

A. Maier, V. Christlein, K. Breininger, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle,
K. Packhäuser, M. Zinnen

Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 24, 2023



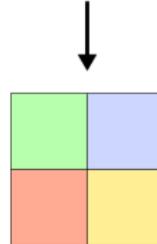
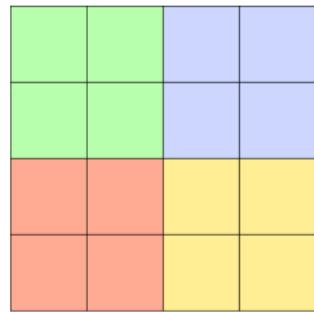
Pooling Layers

Idea behind Pooling Layers

adder info

- Fuses information of input across spatial locations
- Decreases number of parameters
- Reduces computational costs and overfitting

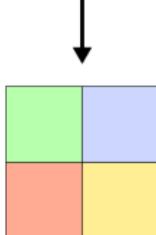
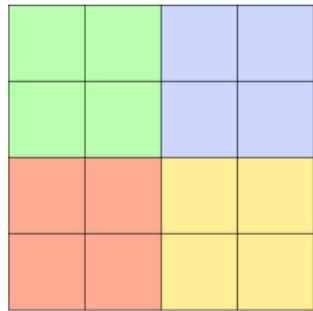
Introduces → Hierarchy



Idea behind Pooling Layers

[C.S]

- Fuses information of input across spatial locations
- Decreases number of parameters
- Reduces computational costs and overfitting
- Assumptions:
 - Features are hierarchically structured
 - Creates “summaries” of regions
 - Provides translational invariance
 - Exact location of a feature is not important



Max Pooling – Forward Pass

- Propagate maximum value in a neighborhood to next layer
- Typical choices: 2×2 or 3×3 neighborhood
- “Stride” of pooling usually equals the neighborhood size
- Maximum propagation adds additional non-linearity → Yes

Max pooling concept. Note that usually a stride > 1 is used for pooling.

Max Pooling – Backward Pass



Max Pooling – Backward Pass



- Only one value contributes to error
- Error is propagated only along the path of the maximum value

rest all get 0

Average Pooling

- Propagate average of the neighborhood
- Does not consistently perform better than max pooling
- Backward pass: Error is shared to equal parts



Additional Pooling Strategies

- Fractional max pooling
- L_p pooling
- Stochastic pooling
- Spacial pyramid pooling
- Generalized pooling
- ...



Many strategies

Alternative: Strided Convolution

- Historically, max pooling was the most frequently used pooling strategies due to additional non-linearity → Max \Rightarrow Non-linearity.
- More recently, convolution with stride $s > 1$ has become more common
 - Allows for trainable downsampling strategy

→ since $s > 1$

Don't need to
encode
max pooling

Conv pooling
= strided Conv

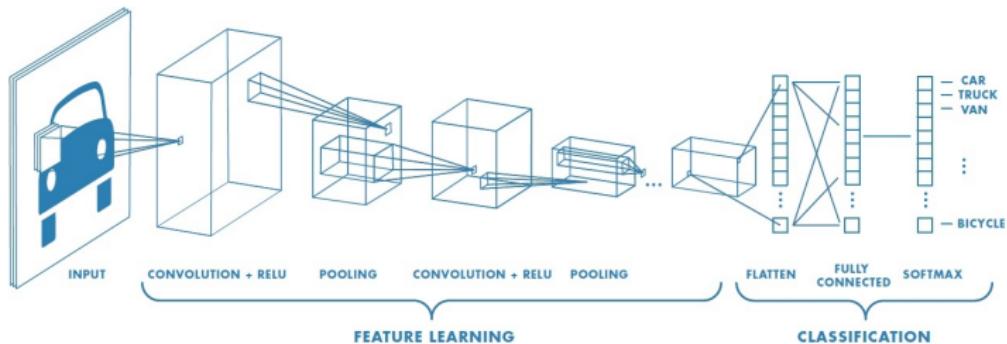
CRP

SAM

ERP

CAP

Recap: Convolutional Neural Networks - Architecture

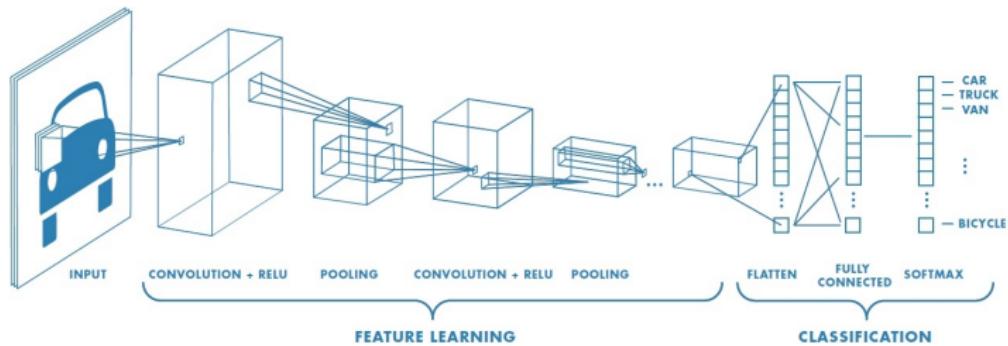


Four essential building blocks:

- Convolutional layer: Feature extraction
- Activation function: Nonlinearity
- Pooling layer: Compress and aggregate information, save parameters
- Last layer: Fully-connected for classification

Source: <https://de.mathworks.com/discovery/convolutional-neural-network.html>

Recap: Convolutional Neural Networks - Architecture

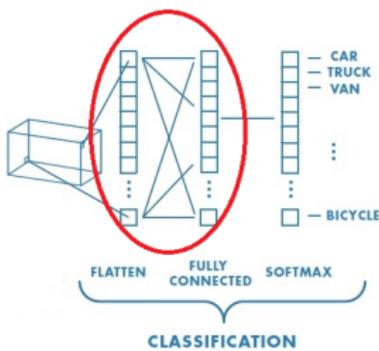


Four Three essential building blocks:

- Convolutional layer: Feature extraction
- Activation function: Nonlinearity
- Pooling layer: Compress and aggregate information, save parameters
- Last layer: Fully connected for classification **We can replace this layer!**

Source: <https://de.mathworks.com/discovery/convolutional-neural-network.html>

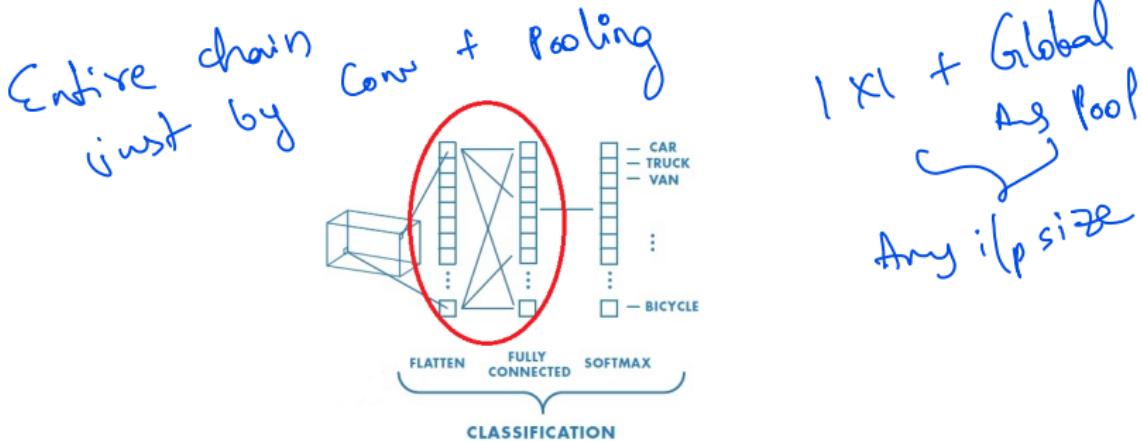
Replacing the Fully Connected Layer



- Conv and pooling layers generate better representation → better features
- Fully connected layers for classification

Flatten + 1×1 = full

Replacing the Fully Connected Layer



- Conv and pooling layers generate better representation → better features
- Fully connected layers for classification
- **Alternatively and equivalently:** Use flatten & 1×1 convolution or $N \times N$ convolution
Any size.
- Enables **arbitrary input sizes** in combination with **global average pooling!**

Inception model

- Szegedy et al. (2014): Going Deeper With Convolutions [8]
 - Very influential publication: > 17000 citations (Nov. '19)
 - Won ImageNet Large-Scale Visual Recognition Challenge 2014
 - GoogLeNet as one incarnation
 - Inspired by Network in Network [4]

Inception model

- Szegedy et al. (2014): Going Deeper With Convolutions [8]
 - Very influential publication: > 17000 citations (Nov. '19)
 - Won ImageNet Large-Scale Visual Recognition Challenge 2014
 - GoogLeNet as one incarnation
 - Inspired by Network in Network [4]
- Self-stated motto:

Inception model

Dream Deep

Going Deeper

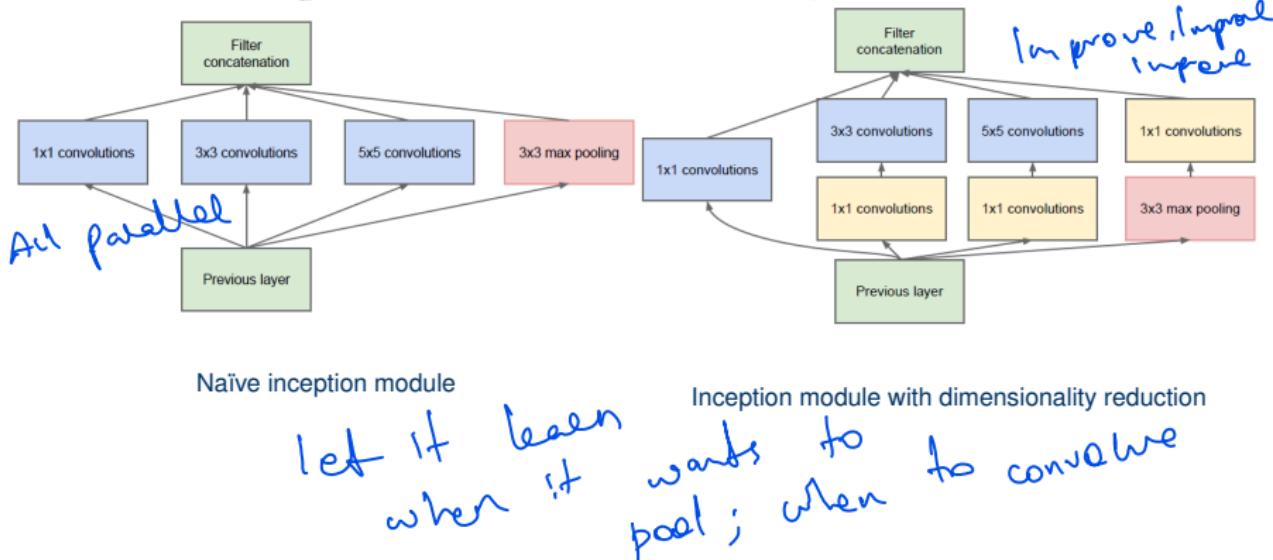
- Szegedy et al. (2014): Going Deeper With Convolutions [8]
 - Very influential publication: > 17000 citations (Nov. '19)
 - Won ImageNet Large-Scale Visual Recognition Challenge 2014
 - GoogLeNet as one incarnation
 - Inspired by Network in Network [4]
- Self-stated motto:



Source: <http://knowyourmeme.com/photos/531557-we-need-to-go-deeper>

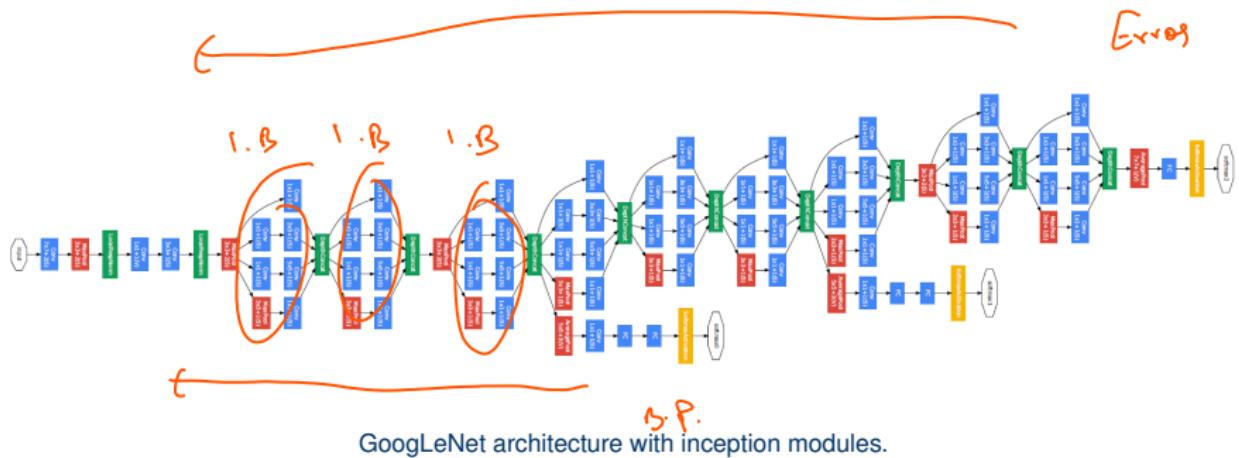
Inception model

- **Idea:** Why use only one type of filter in one layer?
Why not combine different neighborhoods/pooling/etc.?
- Construction of 'inception modules' that are stacked to form a large network.



Source: [8]

Inception model – GoogLeNet



Pretty deep

Source: [8]

NEXT TIME
ON DEEP LEARNING

Coming up

- How to prevent networks from just memorizing the training data?
- Is there a way to force features to be independent?
- How can we make sure our network also recognizes cats in different poses?
- Can we fix the covariate shift problem?

Changes in distribution of
train & test data

Poor Generalization.

Comprehensive Questions

- Name five activation functions.
[C.S] tanh, sign, ReLU, sigmoid
- Discuss those 5 activation functions.
- What is the zero-centering problem? → Action, internal covariate shift
- Why does ReLU as activation function perform much better than sigmoid/tanh in a large number of tasks? → Squeezing NO; hierarchy, receptive field
- Why are convolutional networks well suited for image and audio processing?
- Write down a mathematical description of strided convolution.
- What is the connection between 1×1 convolutions and fully connected layers?
f.l.a + 1x1 = f.c
- How would you implement a classifier which operates on image patches?
- What is a pooling layer? [P.S] C, A, P, Insufficient Data
- Why do we use pooling layers?
- On what data would CNNs probably perform bad? → D.S.W.F.,

Further Reading

- [Link](#) - [3] for a paper about Self Normalizing Networks
- [Link](#) - [4] for a creative Network in Network paper
- [Link](#) - [6] for details on learned activation functions
- [Link](#) - [8] if everything so far was not deep enough for you

Questions?



FAU

FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

References



References I

- [1] I. J. Goodfellow, D. Warde-Farley, M. Mirza, et al. "Maxout Networks". In: [ArXiv e-prints](#) (Feb. 2013). arXiv: 1302.4389 [stat.ML].
- [2] Kaiming He, Xiangyu Zhang, Shaoqing Ren, et al. "Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification". In: [CoRR](#) abs/1502.01852 (2015). arXiv: 1502.01852.
- [3] Günter Klambauer, Thomas Unterthiner, Andreas Mayr, et al. "Self-Normalizing Neural Networks". In: [Advances in Neural Information Processing Systems \(NIPS\)](#). Vol. abs/1706.02515. 2017. arXiv: 1706.02515.
- [4] Min Lin, Qiang Chen, and Shuicheng Yan. "Network In Network". In: [CoRR](#) abs/1312.4400 (2013). arXiv: 1312.4400.
- [5] Andrew L. Maas, Awni Y. Hannun, and Andrew Y. Ng. "Rectifier Nonlinearities Improve Neural Network Acoustic Models". In: [Proc. ICML](#). Vol. 30. 1. 2013.

References II

- [6] Prajit Ramachandran, Barret Zoph, and Quoc V. Le. "Searching for Activation Functions". In: [CoRR abs/1710.05941](#) (2017). arXiv: 1710.05941.
- [7] Stefan Elfwing, Eiji Uchibe, and Kenji Doya. "Sigmoid-weighted linear units for neural network function approximation in reinforcement learning". In: [arXiv preprint arXiv:1702.03118](#) (2017).
- [8] Christian Szegedy, Wei Liu, Yangqing Jia, et al. "Going Deeper with Convolutions". In: [CoRR abs/1409.4842](#) (2014). arXiv: 1409.4842.